

---

## **VE482 - Lab 9**

Introduction to Operating Systems

Weili Shi 519370910011

Kexuan Huang 518370910126

December 3, 2021



### 3. Tasks

#### What needs to be returned by read and write file operations for a character device?

According to *The Linux Kernel Documentation*,

The value returned by read or write can be:

- the number of bytes transferred; if the returned value is less than the size parameter (the number of bytes requested), then it means that a partial transfer was made. Most of the time, the user-space app calls the system call (read or write) function until the required data number is transferred.
- 0 to mark the end of the file in the case of read ; if write returns the value 0 then it means that no byte has been written and that no error has occurred; In this case, the user-space application retries the write call.
- a negative value indicating an error code.

#### How are exactly those major and minor numbers working? You vaguely remember that you can display them using `ls -l /dev`.

According to *The Linux Kernel Documentation*,

The identifier consists of two parts: major and minor.

- The first part identifies the device type (IDE disk, SCSI disk, serial port, etc.)
- The second one identifies the device (first disk, second serial port, etc.).

Most times, the major identifies the driver, while the minor identifies each physical device served by the driver. In general, a driver will have a major associate and will be responsible for all minors associated with that major.

To create a device type file, use the `mknod` command; the command receives the type (block or character), major and minor of the device (`mknod name type major minor`).

Thus, if you want to create a character device named `mycdev` with the major 42 and minor 0, use the command:

```
1 # mknod /dev/mycdev c 42 0
```

## Knowing the major number and minor numbers of a device, how to add a character device to /dev?

With `cdev_add()` and number of major and minor, we can add a character device to `/dev`:

```
1 int cdev_add(struct cdev *dev, dev_t num, unsigned int count);
```

According to *The Linux Kernel Documentation*, the devices can be registered and initialized by:

```
1 #include <linux/fs.h>
2 #include <linux/cdev.h>
3
4 #define MY_MAJOR      42
5 #define MY_MAX_MINORS  5
6
7 struct my_device_data {
8     struct cdev cdev;
9     /* my data starts here */
10    //...
11 };
12
13 struct my_device_data devs[MY_MAX_MINORS];
14
15 const struct file_operations my_fops = {
16     .owner = THIS_MODULE,
17     .open = my_open,
18     .read = my_read,
19     .write = my_write,
20     .release = my_release,
21     .unlocked_ioctl = my_ioctl
22 };
23
24 int init_module(void)
25 {
26     int i, err;
27     err = register_chrdev_region(MKDEV(MY_MAJOR, 0), MY_MAX_MINORS
28                                , "my_device_driver");
29     if (err != 0) {
30         /* report error */
31         return err;
32     }
33     for(i = 0; i < MY_MAX_MINORS; i++) {
34         /* initialize devs[i] fields */
35         cdev_init(&devs[i].cdev, &my_fops);
36         cdev_add(&devs[i].cdev, MKDEV(MY_MAJOR, i), 1);
37     }
38     return 0;
39 }
```

**Where are the following terms located in linux source code?**

According to results in *Bootlin Elixir Cross Referencer*:

- `module_init` `include/linux/module.h`
- `module_exit` `include/linux/module.h`
- `printk` `include/linux/printk.h`
- `container_of` `include/linux/kernel.h`
- `dev_t` `include/linux/types.h`
- `MAJOR` `include/linux/kdev_t.h`
- `MINOR` `include/linux/kdev_t.h`
- `MKDEV` `include/linux/kdev_t.h`
- `alloc_chrdev_region` `include/linux/fs.h`
- `module_param` `include/linux/moduleparam.h`
- `cdev_init` `include/linux/cdev.h`
- `cdev_add` `include/linux/cdev.h`
- `cdev_del` `include/linux/cdev.h`
- `THIS_MODULE` `include/linux/export.h`

**How to generate random numbers when working inside the Linux kernel? You think that a while back you read something about getting the current time.**

```
void get_random_bytes(void *buf, int nbytes);
```

Returns the requested number of random bytes and stores them in a buffer.

## How to define and specify module options?

According to *Linux manual page*,

The `module_init()` macro defines which function is to be called at module insertion time (if the file is compiled as a module), or at boot time (if the file is not compiled as a module the `module_init()` macro becomes equivalent to `__initcall()`, which through linker magic ensures that the function is called on boot.)

- To load a module by filename:  
# `insmod <filename> [args]`
- To pass a parameter to a kernel module,  
# `modprobe <module_name> [parameter=value]`